



北京邮电大学软件学院

School Of software Engineering Of BUPT

厚德 博学 敬业 乐群



Operating Systems

Lecture 6 CPU Scheduling

Jinpengchen

Email: jpchen@bupt.edu.cn



Catalog Description

- ⊕ Basic Concepts
- ⊕ Scheduling Criteria
- ⊕ Scheduling Algorithms
- ⊕ Algorithm Evaluation



Chapter Objectives

- ✿ To introduce CPU scheduling, which is the basis for multiprogrammed operating systems
- ✿ To describe various CPU-scheduling algorithms
- ✿ To discuss evaluation criteria for selecting a CPU-scheduling algorithm for a particular system.



Basic Concepts

- ❁ Scheduling is a fundamental OS function.
 - ❏ Almost all computer resources are scheduled before use.
 - ❏ CPU scheduling is the basis of multiprogrammed OSes.
- ❁ Objective of multiprogramming
 - ❏ Maximum CPU utilization



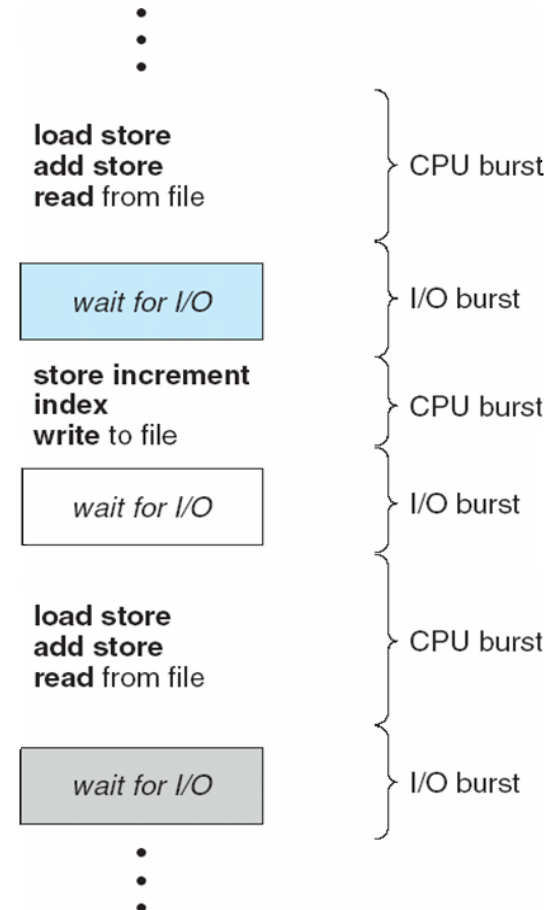
Basic Concepts

❖ CPU-I/O Burst Cycle

- ❖ A property of process :

CPU-I/O Burst Cycle

- ❖ Process execution consists of a cycle of CPU execution and I/O wait
- ❖ Alternating Sequence of CPU and I/O Bursts
- ❖ Begin and end with a CPU burst
- ❖ Process execution
= n (CPU execution + I/O wait)
+ CPU execution





Basic Concepts

✚ CPU-I/O Burst Cycle

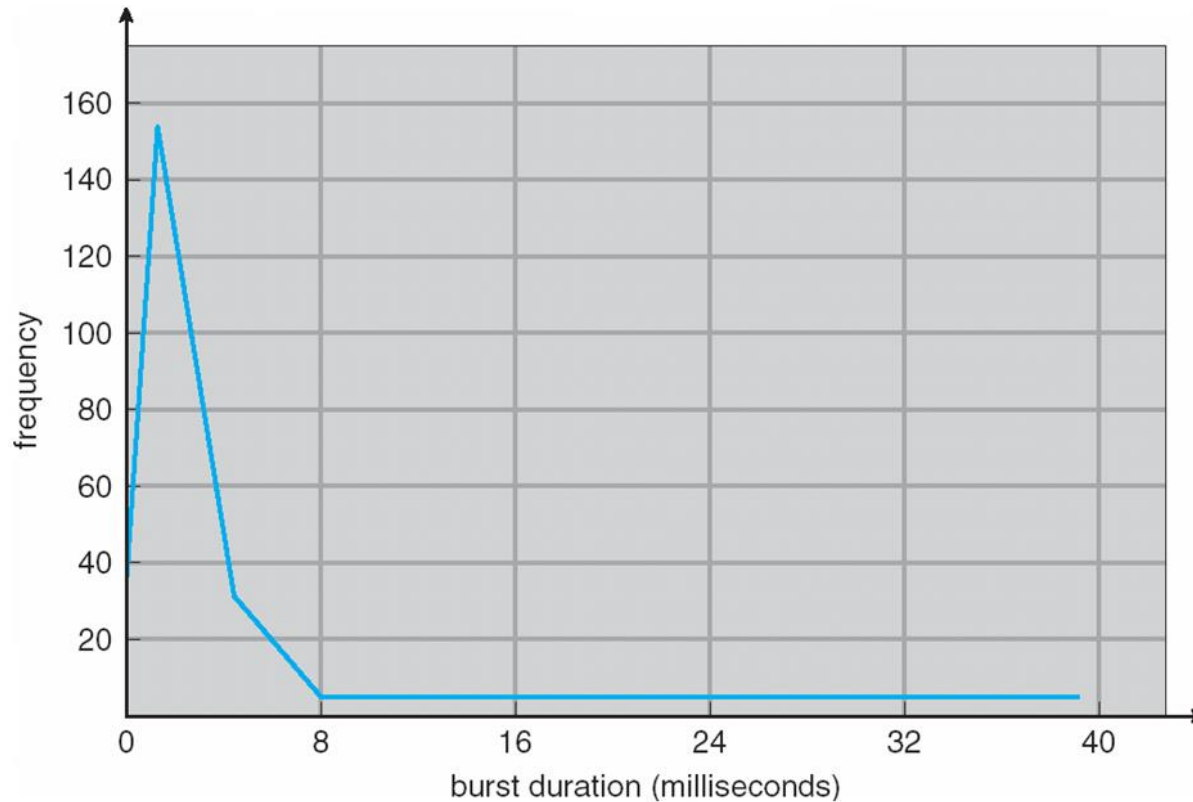
✚ CPU burst distribution

- ✓ CPU burst curve - exponential or hyper exponential
- ✓ CPU-bound program
- ✓ I/O-bound program



Basic Concepts

✚ Histogram of CPU-burst Times





Basic Concepts

❖ CPU Scheduler (Short-term Scheduler)

- ❖ selects a process from the processes in memory that are ready to execute and **allocates** the CPU to that process
- ❖ CPU scheduling decisions may take place when a process:
 - ✓ Switches from **running** to **waiting** state
 - ✓ Switches from **running** to **ready** state
 - ✓ Switches from **waiting** to **ready** state
 - ✓ **Terminates**
- ❖ Scheduling under 1 and 4 is **nonpreemptive**
- ❖ All other scheduling is **preemptive**



Basic Concepts

❁ Dispatcher

- ❁ Dispatcher module gives control of the CPU to the process selected by the short-term scheduler; this involves:
 - ✓ switching context
 - ✓ switching to user mode
 - ✓ jumping to the proper location in the user program to restart that program
- ❁ Dispatch latency - the time it takes for the dispatcher to stop one process and start another running
 - ✓ SHOULD be as fast as possible



Catalog Description

- ⊕ Basic Concepts
- ⊕ Scheduling Criteria
- ⊕ Scheduling Algorithms
- ⊕ Algorithm Evaluation



Scheduling Criteria

- ⊕ CPU utilization (CPU 利用率)
 - keep the CPU as busy as possible
- ⊕ Throughput (吞吐率)
- ⊕ Turnaround time (周转时间)
- ⊕ Waiting time (等待时间)
- ⊕ Response time (响应时间)



Scheduling Criteria

❁ Throughput

- ❏ the number of processes that are completed per time unit
 - ✓ different from one process set to another process set
 - ✓ for long processes: may be 1 process per hour
 - ✓ for short transactions: may be 10 processes per second



Scheduling Criteria

✚ Turnaround time

- ✚ the amount of time to execute a particular process
 - ✓ from the time of submission of a process to the time of completion
 - ✓ = the sum of the periods spent waiting to get into memory, waiting in the ready queue, executing on the CPU, and doing I/O.



Scheduling Criteria

⊕ Waiting time

- ⊞ the amount of time a process has been waiting in the ready queue
- ⊞ = the sum of the periods spent waiting in the ready queue

⊕ Response time

- ⊞ amount of time it takes from when a request was submitted until the first response is produced, not the time it takes to output the response
- ✓ for time-sharing environment



Scheduling Algorithm Optimization Criteria

- ⊕ Max CPU utilization
- ⊕ Max throughput
- ⊕ Min turnaround time
- ⊕ Min waiting time
- ⊕ Min response time



Catalog Description

- ⊕ Basic Concepts
- ⊕ Scheduling Criteria
- ⊕ Scheduling Algorithms
- ⊕ Algorithm Evaluation



Scheduling Algorithms

- ✿ FCFS (先来先服务) Scheduling
- ✿ SJF (短作业优先) Scheduling
- ✿ Priority Scheduling (优先级)
- ✿ Round Robin (时间片轮转) Scheduling
- ✿ Multilevel Queue (多级队列) Scheduling
- ✿ Multilevel Feedback Queue (多级反馈队列) Scheduling



Scheduling Algorithms

✿ FCFS (先来先服务) Scheduling

- ✦ Nonpreemptive (非抢占)

- ✦ Implementation

 - ✓ Normal Queue: FIFO Queue

 - ✓ Ordered by request time

 - ✓ Linked list

 - ✓ Insert: linked to the tail of the queue

 - ✓ Scheduling: removed from the head of the queue



Scheduling Algorithms

- Suppose that the processes arrive in the order: P1 , P2 , P3

<u>Process</u>	<u>Burst Time(ms)</u>
P_1	24
P_2	3
P_3	3

- The Gantt Chart for the schedule is:



- Waiting time for P1 = 0; P2 = 24; P3 = 27
- Average waiting time: $(0 + 24 + 27) / 3 = 17$



Scheduling Algorithms

- ✿ Suppose that the processes arrive in the **order**: P2, P3, P1
- ✿ The **Gantt Chart** for the schedule is:



- ✿ **Waiting time** for P1 = 6; P2 = 0; P3 = 3
- ✿ **Average waiting time**: $(6 + 0 + 3)/3 = 3$
- ✿ Much better than previous case



Scheduling Algorithms

- ✿ Convoy effect (护航效应; 护卫效应)
 - ✦ all the other processes wait for the one big process to get off the CPU
 - ✦ \equiv short process behind long process
- ✿ Example situation:
 - ✦ one CPU-bound process
 - ✦ many I/O-bound processes



Scheduling Algorithms

✿ Shortest-Job-First (SJF) Scheduling

- ✦ Associate with each process the length of its next CPU burst.
- ✦ Schedule the process with **the shortest next CPU burst**.




Scheduling Algorithms

❖ SJF scheduling example

Process	Burst Time (ms)
P1	6
P2	8
P3	7
P4	3

❖ The Gantt chart for the schedule is:



- ❖ Waiting time for P1= 3; P2 = 16; P3 = 9; P4 = 0
- ❖ Average waiting time: $(3 + 16 + 9 + 0)/4 = 7$
- ❖ If FCFS, average waiting time: $(0 + 6 + 14 + 21)/4 = 10.25$



Scheduling Algorithms

Shortest-Job-First (SJF) Scheduling

❑ Two schemes:

✓ **nonpreemptive**

- once the CPU is given to the process, it cannot be preempted until it completes its CPU burst

✓ **preemptive**

- if a new process arrives with CPU burst length less than the remaining time of current executing process, preempt. This scheme is known as the **Shortest-Remaining-Time-First (SRTF)**

❑ SJF is **optimal** (最优的)

- gives the minimum average waiting time (最小平均等待时间) for a given set of processes

✓ The **difficulty** is knowing the length of the next CPU request




Scheduling Algorithms

Example of Non-Preemptive SJF

Process	Arrival Time	Burst Time(ms)
P1	0.0	7
P2	2.0	4
P3	4.0	1
P4	5.0	4

The Gantt chart for SJF (non-preemptive)



Average waiting time: $(0 + 6 + 3 + 7)/4 = 4$




Scheduling Algorithms

Example of Preemptive SJF

Process	Arrival Time	Burst Time
P1	0.0	7
P2	2.0	4
P3	4.0	1
P4	5.0	4

The Gantt chart for SJF (preemptive)



Average waiting time: $((11-2) + (5-4) + 0 + (7-5)) / 4 = 3$



Scheduling Algorithms

❁ Determining Length of Next CPU Burst

- ❁ can only **estimate** the length
- ❁ can be done by using the length of previous CPU bursts, using **exponential averaging** (指数平均)
 - ✓ t_n = actual length of n^{th} CPU burst
 - ✓ τ_{n+1} = predicted value for the next CPU burst
 - ✓ α , $0 \leq \alpha \leq 1$
 - ✓ Define: $\tau_{n+1} = \alpha t_n + (1 - \alpha) \tau_n$



Scheduling Algorithms

Examples of Exponential Averaging

■ $\alpha = 0$

- $\tau_{n+1} = \tau_n$
- Recent history does not count

■ $\alpha = 1$

- $\tau_{n+1} = \alpha t_n$
- Only the actual last CPU burst counts

■ If we expand the formula, we get:

$$\begin{aligned}\tau_{n+1} = & \alpha t_n + (1 - \alpha)\alpha t_{n-1} + \dots \\ & + (1 - \alpha)^j \alpha t_{n-j} + \dots \\ & + (1 - \alpha)^{n+1} \tau_0\end{aligned}$$


- ### ■ Since both α and $(1 - \alpha)$ are less than or equal to 1, each successive term has less weight than its predecessor





Scheduling Algorithms

- Prediction of the Length of the Next CPU Burst
- Example: $\alpha = 1/2$, $\tau_0 = 10$



CPU burst (t_i)	6	4	6	4	13	13	13	...	
"guess" (τ_i)	10	8	6	6	5	9	11	12	...



Scheduling Algorithms

✚ Priority Scheduling

- ✚ A priority number(优先数) is associated with each process
- ✚ The CPU is allocated to the process with the highest priority
- ✚ Priority scheduling can be:
 - ✓ Preemptive VS. Nonpreemptive
- ✚ SJF is a special case of general priority scheduling where priority is the (predicted) next CPU burst time
- ✚ Problem – **Starvation** – low priority processes may never execute
- ✚ Solution – **Aging** – as time progresses, the priority of the process can be increased




Scheduling Algorithms

Example of Non-Preemptive Priority Scheduling

Process	Burst Time (ms)	Priority
P1	10.0	3
P2	1.0	1
P3	2.0	4
P4	1.0	5
P5	5.0	2

The Gantt chart for Non-Preemptive Priority



Average waiting time: $(6 + 0 + 16 + 18 + 1) / 5 = 8.2$



Scheduling Algorithms

Round Robin (时间片轮转, RR) Scheduling

- Each process gets a small unit of CPU time (time quantum), usually 10–100 milliseconds. After this time has elapsed, the process is preempted and added to the end of the ready queue.
- If there are n processes in the ready queue and the time quantum is q , then each process gets $1/n$ of the CPU time in chunks of at most q time units at once.
- No process waits more than $(n-1)q$ time units.
- Performance**
 - ✓ q large – FCFS
 - ✓ q small – q must be large with respect to context switch, otherwise **overhead** is too high




Scheduling Algorithms

Example of RR with Time Quantum = 20

Process	Burst Time
P1	53
P2	17
P3	68
P4	24

The Gantt chart is:




Typically, higher average turnaround time than SJF, but better response



Scheduling Algorithms

Time Quantum and Context Switch Time

- The effect of context switching on the performance of RR scheduling




- typically the context-switch time is a small fraction of the time quantum
 - usually: time quantum: 10 ~ 100ms & context switch time: 10 μ s



Scheduling Algorithms

Turnaround Time Varies With the Time Quantum



process	time
P_1	6
P_2	3
P_3	1
P_4	7



Scheduling Algorithms


❖ Multilevel Queue (多级队列)

- ❖ Ready queue is partitioned into separate queues:
 - foreground (interactive)
 - background (batch)
- ❖ Each queue has its own scheduling algorithm
 - ✓ foreground - RR
 - ✓ background - FCFS
- ❖ Scheduling must be done between the queues
 - ✓ Fixed priority scheduling;
 - Example: serve all from foreground then from background).
 - Possibility of starvation.
 - ✓ Time slice - each queue gets a certain amount of CPU time which it can schedule amongst its processes; i.e.,
 - 80% to foreground in RR
 - 20% to background in FCFS



Scheduling Algorithms

❁ Multilevel Queue Scheduling





Scheduling Algorithms

❖ Multilevel Feedback Queue (多级反馈队列) Scheduling

- ❖ A process can **move** between the various queues; aging can be implemented in this way
- ❖ Multilevel-feedback-queue(多级反馈队列) scheduler defined by the following parameters
 - ✓ **number of queues**
 - ✓ **scheduling algorithms** for each queue
 - ✓ method used to determine when to **upgrade** a process
 - ✓ method used to determine when to **demote** a process
 - ✓ method used to determine which queue a process will enter when that process needs service



Scheduling Algorithms

✚ Example of Multilevel Feedback Queue

✚ Three queues:

- ✓ Q0 - RR with time quantum 8 milliseconds
- ✓ Q1 - RR time quantum 16 milliseconds
- ✓ Q2 - FCFS


✚ Scheduling

- ✓ A new job enters queue Q0 which is served FCFS. When it gains CPU, job receives 8 milliseconds. If it does not finish in 8 milliseconds, job is moved to queue Q1.
- ✓ At Q1 job is again served FCFS and receives 16 additional milliseconds. If it still does not complete, it is preempted and moved to queue Q2.



Scheduling Algorithms

Example of Multilevel Feedback Queue





Catalog Description

- ⊕ Basic Concepts
- ⊕ Scheduling Criteria
- ⊕ Scheduling Algorithms
- ⊕ Algorithm Evaluation



Algorithm Evaluation

Algorithm Evaluation

- ❖ How do we select a CPU scheduling algorithm for a particular system?
 - ✓ firstly, which criteria? What is the relative importance of these measures
 - ✓ then, evaluate the algorithms
 - Deterministic Modeling(确定性建模)
 - Queueing Models(排队模型)
 - Simulations(模拟)
 - Implementation Scheduling



Algorithm Evaluation

- ❖ Deterministic Modeling(确定性建模)
 - ❖ Analytic evaluation(分析评估法): One major class of evaluation methods
 - ✓ uses the given algorithm and the system workload to produce a formula or number that evaluates the performance of the algorithm for that workload.
 - ❖ Deterministic modeling(确定性建模) – takes a particular predetermined workload and defines the performance of each algorithm for that workload



Algorithm Evaluation

❖ Queuing Models (排队模型)

- ❖ Usually, two distributions can be measured and then approximated or simply estimated
 - ✓ the distribution of CPU and I/O bursts
 - ✓ the arrival-time distribution
- ❖ Queueing-network analysis (排队网络分析)
 - ✓ Computer System: a network of servers, each server has a queue of waiting processes
 - CPU: ready queue;
 - I/O: device queues (\equiv waiting queue)
 - ✓ Given arriving rates and service rates
 \Rightarrow utilization,
average queue length, average wait time, ...



Algorithm Evaluation

❖ Simulations (模拟)

❖ Running simulations involves **programming a model of the computer system**.

- ✓ Software data structures represent the major components
 - a clock
 - the system state is modified to reflect the activities of the devices, the processes and the scheduler.
- ✓ finally, the statistics are gathered

❖ How to generate the data to drive the simulation?

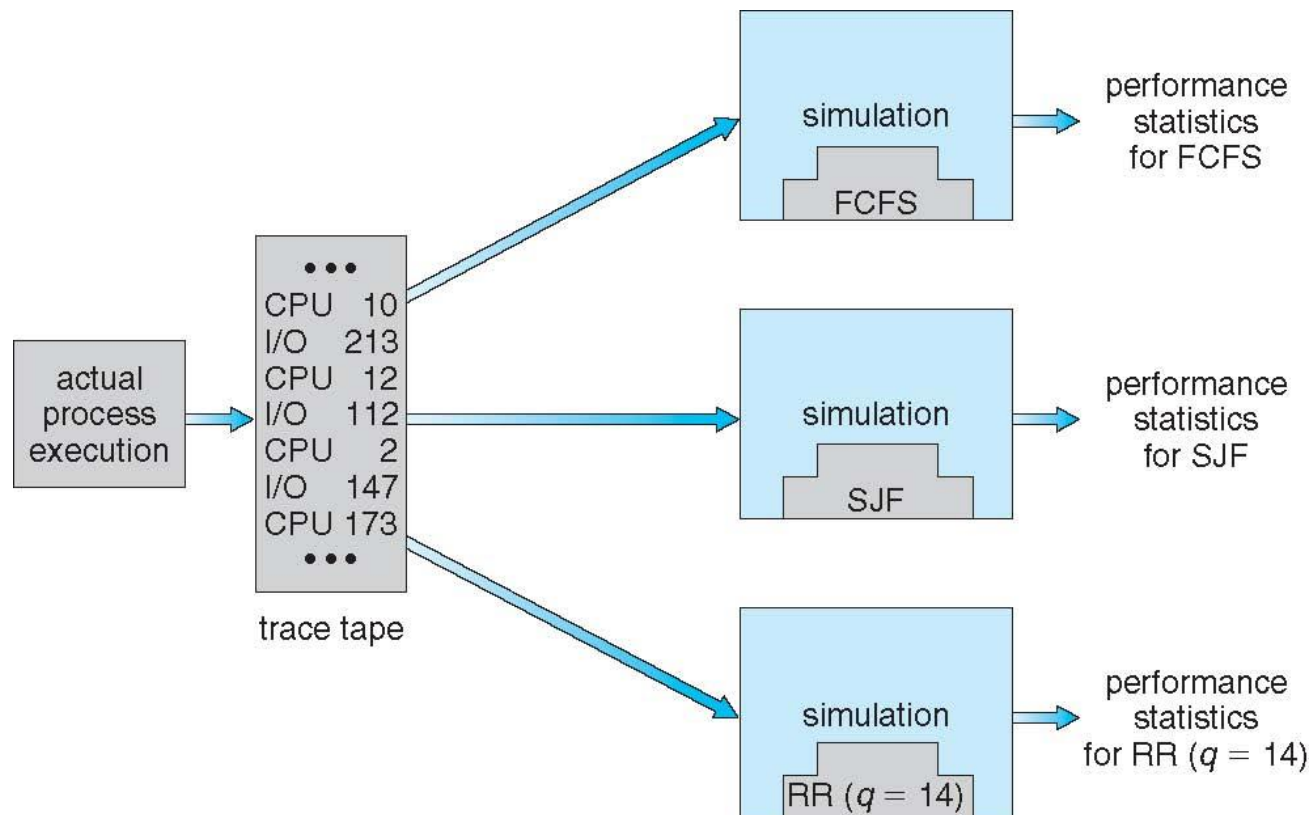
- ✓ **distribution-driven simulation**
 - random-number generator, according to probability distributions, to generate processes, CPU burst times, arrivals, departures, ...
 - the distributions can be defined mathematically (uniform, exponential, Poisson) or empirically
 - may be inaccurate



Algorithm Evaluation

❁ Evaluation of CPU schedulers by Simulation

❁ trace tapes (跟踪磁带)





Algorithm Evaluation

❖ Implementation

- ❖ This approach puts the actual algorithm in the real system for evaluation under real operating conditions
- ❖ the main difficulty: high cost



End of Chapter 6